

# Generating new MNIST digits with GAN

Source: [https://github.com/Zackory/Keras-MNIST-GAN/blob/master/mnist\\_gan.py](https://github.com/Zackory/Keras-MNIST-GAN/blob/master/mnist_gan.py)

```
In [2]: ▶ import os
import numpy as np
from tqdm import tqdm
import matplotlib.pyplot as plt

from tensorflow import keras
from keras.layers import Input
from keras.models import Model, Sequential
from keras.layers.core import Reshape, Dense, Dropout, Flatten
from keras.layers.advanced_activations import LeakyReLU
from keras.layers.convolutional import Convolution2D, UpSampling2D
from keras.layers.normalization import BatchNormalization
from keras.datasets import mnist
from keras.optimizers import Adam
from keras import backend as K
from keras import initializers
```

```
In [3]: ▶ # Deterministic output.
# Tired of seeing the same results every time? Remove the line below.
np.random.seed(1000)

# The results are a little better when the dimensionality of the random vector is only 10.
# The dimensionality has been left at 100 for consistency with other GAN implementations.
randomDim = 100
```

## Step 1: Load the dataset

```
In [4]: ▶ # Load MNIST data
(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = (X_train.astype(np.float32) - 127.5)/127.5
X_train = X_train.reshape(60000, 784)
```

```
In [5]: ▶ # Optimizer
adam = Adam(lr=0.0002, beta_1=0.5)
```

## Step 2: Build the generator model

```
In [8]: generator = Sequential(name = 'generator')
generator.add(Dense(256, input_dim=randomDim, kernel_initializer=initializers.RandomNormal(stddev=0.02)))
generator.add(LeakyReLU(0.2))
generator.add(Dense(512))
generator.add(LeakyReLU(0.2))
generator.add(Dense(1024))
generator.add(LeakyReLU(0.2))
generator.add(Dense(784, activation='tanh'))
generator.compile(loss='binary_crossentropy', optimizer=adam)
generator.summary()
```

Model: "generator"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 256)	25856
leaky_re_lu_6 (LeakyReLU)	(None, 256)	0
dense_9 (Dense)	(None, 512)	131584
leaky_re_lu_7 (LeakyReLU)	(None, 512)	0
dense_10 (Dense)	(None, 1024)	525312
leaky_re_lu_8 (LeakyReLU)	(None, 1024)	0
dense_11 (Dense)	(None, 784)	803600
Total params: 1,486,352		
Trainable params: 1,486,352		
Non-trainable params: 0		

### Step 3: Build the discriminator model

```
In [9]: discriminator = Sequential(name = 'discriminator')
discriminator.add(Dense(1024, input_dim=784, kernel_initializer=initializers.RandomNormal(stddev=0.02)))
discriminator.add(LeakyReLU(0.2))
discriminator.add(Dropout(0.3))
discriminator.add(Dense(512))
discriminator.add(LeakyReLU(0.2))
discriminator.add(Dropout(0.3))
discriminator.add(Dense(256))
discriminator.add(LeakyReLU(0.2))
discriminator.add(Dropout(0.3))
discriminator.add(Dense(1, activation='sigmoid'))
discriminator.compile(loss='binary_crossentropy', optimizer=adam)
discriminator.summary()
```

Model: "discriminator"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 1024)	803840
leaky_re_lu_9 (LeakyReLU)	(None, 1024)	0
dropout_3 (Dropout)	(None, 1024)	0
dense_13 (Dense)	(None, 512)	524800
leaky_re_lu_10 (LeakyReLU)	(None, 512)	0
dropout_4 (Dropout)	(None, 512)	0
dense_14 (Dense)	(None, 256)	131328
leaky_re_lu_11 (LeakyReLU)	(None, 256)	0
dropout_5 (Dropout)	(None, 256)	0
dense_15 (Dense)	(None, 1)	257
Total params: 1,460,225		
Trainable params: 1,460,225		
Non-trainable params: 0		

Step 4: Build GAN model = generator + discriminator

```
In [10]: # Combined network
discriminator.trainable = False
ganInput = Input(shape=(randomDim,))
x = generator(ganInput)
ganOutput = discriminator(x)
gan = Model(inputs=ganInput, outputs=ganOutput)
gan.compile(loss='binary_crossentropy', optimizer=adam)
gan.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 100)]	0
-----		
generator (Sequential)	(None, 784)	1486352
-----		
discriminator (Sequential)	(None, 1)	1460225
=====		
Total params: 2,946,577		
Trainable params: 1,486,352		
Non-trainable params: 1,460,225		

Step 5: Prepare functions to produce and plot the results

```

In [25]: dLosses = []
         gLosses = []

# Plot the Loss from each batch
def plotLoss(epoch):
    plt.figure(figsize=(10, 8))
    plt.plot(dLosses, label='Discriminative loss')
    plt.plot(gLosses, label='Generative loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.savefig('images/gan_loss_epoch_%d.png' % epoch)

# Create a wall of generated MNIST images
def plotGeneratedImages(epoch, examples=100, dim=(10, 10), figsize=(10, 10)):
    noise = np.random.normal(0, 1, size=[examples, randomDim])
    generatedImages = generator.predict(noise)
    generatedImages = generatedImages.reshape(examples, 28, 28)

    plt.figure(figsize=figsize)
    for i in range(generatedImages.shape[0]):
        plt.subplot(dim[0], dim[1], i+1)
        plt.imshow(generatedImages[i], interpolation='nearest', cmap='gray_r')
        plt.axis('off')
    plt.tight_layout()
    plt.savefig('images/gan_generated_image_epoch_%d.png' % epoch)

# Save the generator and discriminator networks (and weights) for later use
def saveModels(epoch):
    generator.save('models/gan_generator_epoch_%d.h5' % epoch)
    discriminator.save('models/gan_discriminator_epoch_%d.h5' % epoch)

def train(epochs=1, batchSize=128):
    batchCount = X_train.shape[0] / batchSize
    print('Epochs:', epochs)
    print('Batch size:', batchSize)
    print('Batches per epoch:', batchCount)

    for e in range(1, epochs+1):
        print('-'*15, 'Epoch %d' % e, '-'*15)
        for i in tqdm(range(int(batchCount))):
            # Get a random set of input noise and images
            noise = np.random.normal(0, 1, size=[batchSize, randomDim])
            imageBatch = X_train[np.random.randint(0, X_train.shape[0], size=batchSize)]

            # Generate fake MNIST images
            generatedImages = generator.predict(noise)
            # print np.shape(imageBatch), np.shape(generatedImages)
            X = np.concatenate([imageBatch, generatedImages])

            # Labels for generated and real data
            yDis = np.zeros(2*batchSize)
            # One-sided Label smoothing
            yDis[:batchSize] = 0.9

            # Train discriminator
            discriminator.trainable = True
            dloss = discriminator.train_on_batch(X, yDis)

            # Train generator
            noise = np.random.normal(0, 1, size=[batchSize, randomDim])
            yGen = np.ones(batchSize)
            discriminator.trainable = False
            gloss = gan.train_on_batch(noise, yGen)

            # Store loss of most recent batch from this epoch
            dLosses.append(dloss)
            gLosses.append(gloss)

            if e == 1 or e % 20 == 0:
                plotGeneratedImages(e)
                saveModels(e)

# Plot losses from every epoch
plotLoss(e)

```

## Step 6: Fit GAN model & plot results

In [27]: ▶ train(100, 128)

```
100%|██████████| 468/468 [00:40<00:00, 11.46it/s]
  0%|          | 2/468 [00:00<00:40, 11.48it/s]

----- Epoch 99 -----

100%|██████████| 468/468 [00:39<00:00, 11.76it/s]
  0%|          | 2/468 [00:00<00:38, 12.24it/s]

----- Epoch 100 -----

100%|██████████| 468/468 [00:39<00:00, 11.99it/s]
```

In [ ]: ▶