

# Generating new MNIST digits with VAE

Source: <https://www.kaggle.com/code/vincentman0403/vae-with-convolution-on-mnist>

**How to use VAE(Variational Auto-Encoder) to generate different MNIST images? VAE use convolution layers in encoder and decoder. ¶**

```
In [1]: from keras.layers import Dense, Input
from keras.layers import Conv2D, Flatten, Lambda
from keras.layers import Reshape, Conv2DTranspose
from keras.models import Model
from keras.datasets import mnist
from keras.losses import mse, binary_crossentropy
from keras import backend as K

import numpy as np
import matplotlib.pyplot as plt
import argparse
import os
```

## Step 1: Load and explore the dataset

```
In [2]: (x_train, y_train), (x_test, y_test) = mnist.load_data()

image_size = x_train.shape[1]
# Only get some data to train and test
train_len = 10000
test_len = 1000
x_train = np.reshape(x_train[:train_len], [-1, image_size, image_size, 1])
x_test = np.reshape(x_test[:test_len], [-1, image_size, image_size, 1])
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255
y_test = y_test[:test_len]
```

```
In [3]: x_train.shape
```

Out[3]: (10000, 28, 28, 1)

```
In [4]: x_test.shape
```

Out[4]: (1000, 28, 28, 1)


```
In [6]: ▶ #Show MNIST images
x_train_show = x_train.reshape(-1,28,28)
plt.figure(1)
plt.subplot(221)
plt.imshow(x_train_show[20])
print('y_train[20] = ', y_train[20])

plt.subplot(222)
plt.imshow(x_train_show[500])
print('y_train[500] = ', y_train[500])

plt.subplot(223)
plt.imshow(x_train_show[3000])
print('y_train[3000] = ', y_train[3000])

plt.subplot(224)
plt.imshow(x_train_show[9000])
print('y_train[9000] = ', y_train[9000])
plt.show()
```

```
y_train[20] = 4
y_train[500] = 3
y_train[3000] = 9
y_train[9000] = 6
```



## Step 2: Prepare functions to produce and plot the results (before we start)

```
In [7]: ▶ # reparameterization trick: instead of sampling from Q(z|X), sample epsilon = N(0,1)
# then z = z_mean + sqrt(var)*eps
def sampling(args):
    z_mean, z_log_var = args
    batch = K.shape(z_mean)[0]
    dim = K.int_shape(z_mean)[1]
    # by default, random_normal has mean=0 and std=1.0
    epsilon = K.random_normal(shape=(batch, dim))
    return z_mean + K.exp(0.5 * z_log_var) * epsilon
```

```
In [26]: ▶ def plot_results(encoder, decoder, x_test, y_test, batch_size=128):
# display a 2D plot of the digit classes in the latent space
z_mean, _, _ = encoder.predict(x_test,
                                batch_size=batch_size)

plt.figure(figsize=(12, 10))
plt.scatter(z_mean[:, 0], z_mean[:, 1], c=y_test)
plt.colorbar()
plt.xlabel("z[0]")
plt.ylabel("z[1]")
plt.show()

# display a 30x30 2D manifold of digits
n = 30
digit_size = 28
figure = np.zeros((digit_size * n, digit_size * n))
# linearly spaced coordinates corresponding to the 2D plot
# of digit classes in the latent space
grid_x = np.linspace(-4, 4, n)
grid_y = np.linspace(-4, 4, n)[::-1]

for i, yi in enumerate(grid_y):
    for j, xi in enumerate(grid_x):
        z_sample = np.array([[xi, yi]])
        x_decoded = decoder.predict(z_sample)
        digit = x_decoded[0].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              j * digit_size: (j + 1) * digit_size] = digit

plt.figure(figsize=(10, 10))
start_range = digit_size // 2
end_range = n * digit_size + start_range
pixel_range = np.arange(start_range, end_range, digit_size)
sample_range_x = np.round(grid_x, 1)
sample_range_y = np.round(grid_y, 1)
plt.xticks(pixel_range, sample_range_x)
plt.yticks(pixel_range, sample_range_y)
plt.xlabel("z[0]")
plt.ylabel("z[1]")
plt.imshow(figure, cmap='Greys_r')
plt.show()
```

### Step 3: Set up network parameters

```
In [12]: ▶ input_shape = (image_size, image_size, 1)
batch_size = 32
kernel_size = 3
filters = 16
latent_dim = 2
epochs = 30
use_mse = True
load_weights = False
```

### Step 4: Build encoder model

```

In [13]: inputs = Input(shape=input_shape, name='encoder_input')
x = inputs
for i in range(2):
    filters *= 2
    x = Conv2D(filters=filters,
               kernel_size=kernel_size,
               activation='relu',
               strides=2,
               padding='same')(x)

# shape info needed to build decoder model
shape = K.int_shape(x)

# generate latent vector Q(z/X)
x = Flatten()(x)
x = Dense(16, activation='relu')(x)
z_mean = Dense(latent_dim, name='z_mean')(x)
z_log_var = Dense(latent_dim, name='z_log_var')(x)

# use reparameterization trick to push the sampling out as input
# note that "output_shape" isn't necessary with the TensorFlow backend
z = Lambda(sampling, output_shape=(latent_dim,), name='z')([z_mean, z_log_var])

# instantiate encoder model
encoder = Model(inputs, [z_mean, z_log_var, z], name='encoder')
encoder.summary()

```

Model: "encoder"

Layer (type)	Output Shape	Param #	Connected to
encoder_input (InputLayer)	[(None, 28, 28, 1)]	0	
conv2d_2 (Conv2D)	(None, 14, 14, 32)	320	encoder_input[0][0]
conv2d_3 (Conv2D)	(None, 7, 7, 64)	18496	conv2d_2[0][0]
flatten_1 (Flatten)	(None, 3136)	0	conv2d_3[0][0]
dense_1 (Dense)	(None, 16)	50192	flatten_1[0][0]
z_mean (Dense)	(None, 2)	34	dense_1[0][0]
z_log_var (Dense)	(None, 2)	34	dense_1[0][0]
z (Lambda)	(None, 2)	0	z_mean[0][0] z_log_var[0][0]

## Step 5: Build decoder model

```
In [14]: latent_inputs = Input(shape=(latent_dim,), name='z_sampling')
x = Dense(shape[1] * shape[2] * shape[3], activation='relu')(latent_inputs)
x = Reshape((shape[1], shape[2], shape[3]))(x)

# use Conv2DTranspose to reverse the conv layers from the encoder
for i in range(2):
    x = Conv2DTranspose(filters=filters,
                        kernel_size=kernel_size,
                        activation='relu',
                        strides=2,
                        padding='same')(x)

    filters //= 2

outputs = Conv2DTranspose(filters=1,
                          kernel_size=kernel_size,
                          activation='sigmoid',
                          padding='same',
                          name='decoder_output')(x)

# instantiate decoder model
decoder = Model(latent_inputs, outputs, name='decoder')
decoder.summary()
```

Model: "decoder"

Layer (type)	Output Shape	Param #
z_sampling (InputLayer)	[(None, 2)]	0
dense_2 (Dense)	(None, 3136)	9408
reshape (Reshape)	(None, 7, 7, 64)	0
conv2d_transpose (Conv2DTran	(None, 14, 14, 64)	36928
conv2d_transpose_1 (Conv2DTr	(None, 28, 28, 32)	18464
decoder_output (Conv2DTransp	(None, 28, 28, 1)	289
Total params: 65,089		
Trainable params: 65,089		
Non-trainable params: 0		

## Step 6: Build VAE model = encoder + decoder

```
In [15]: outputs = decoder(encoder(inputs)[2])
vae = Model(inputs, outputs, name='vae')
```

## Loss = reconstruction loss + KL loss

```
In [16]: if use_mse:
    reconstruction_loss = mse(K.flatten(inputs), K.flatten(outputs))
else:
    reconstruction_loss = binary_crossentropy(K.flatten(inputs),
                                              K.flatten(outputs))

reconstruction_loss *= image_size * image_size
kl_loss = 1 + z_log_var - K.square(z_mean) - K.exp(z_log_var)
kl_loss = K.sum(kl_loss, axis=-1)
kl_loss *= -0.5
vae_loss = K.mean(reconstruction_loss + kl_loss)
vae.add_loss(vae_loss)
```

## Compile model

```
In [17]: vae.compile(optimizer='rmsprop')
vae.summary()
```

Model: "vae"

Layer (type)	Output Shape	Param #	Connected to
encoder_input (InputLayer)	[(None, 28, 28, 1)]	0	
encoder (Functional)	[(None, 2), (None, 2 69076		encoder_input[0][0]
decoder (Functional)	(None, 28, 28, 1)	65089	encoder[0][2]
conv2d_2 (Conv2D)	(None, 14, 14, 32)	320	encoder_input[0][0]
conv2d_3 (Conv2D)	(None, 7, 7, 64)	18496	conv2d_2[0][0]
flatten_1 (Flatten)	(None, 3136)	0	conv2d_3[0][0]
dense_1 (Dense)	(None, 16)	50192	flatten_1[0][0]
z_log_var (Dense)	(None, 2)	34	dense_1[0][0]
z_mean (Dense)	(None, 2)	34	dense_1[0][0]
tf.reshape_1 (TFOpLambda)	(None,)	0	decoder[0][0]
tf.reshape (TFOpLambda)	(None,)	0	encoder_input[0][0]
tf.__operators__.add (TFOpLambda)	(None, 2)	0	z_log_var[0][0]
tf.math.square (TFOpLambda)	(None, 2)	0	z_mean[0][0]
tf.convert_to_tensor (TFOpLambda)	(None,)	0	tf.reshape_1[0][0]
tf.cast (TFOpLambda)	(None,)	0	tf.reshape[0][0]
tf.math.subtract (TFOpLambda)	(None, 2)	0	tf.__operators__.add[0][0] tf.math.square[0][0]
tf.math.exp (TFOpLambda)	(None, 2)	0	z_log_var[0][0]
tf.math.squared_difference (TFOpLambda)	(None,)	0	tf.convert_to_tensor[0][0] tf.cast[0][0]
tf.math.subtract_1 (TFOpLambda)	(None, 2)	0	tf.math.subtract[0][0] tf.math.exp[0][0]
tf.math.reduce_mean (TFOpLambda)	(None,)	0	tf.math.squared_difference[0][0]
tf.math.reduce_sum (TFOpLambda)	(None,)	0	tf.math.subtract_1[0][0]
tf.math.multiply (TFOpLambda)	(None,)	0	tf.math.reduce_mean[0][0]
tf.math.multiply_1 (TFOpLambda)	(None,)	0	tf.math.reduce_sum[0][0]
tf.__operators__.add_1 (TFOpLambda)	(None,)	0	tf.math.multiply[0][0] tf.math.multiply_1[0][0]
tf.math.reduce_mean_1 (TFOpLambda)	(None,)	0	tf.__operators__.add_1[0][0]
add_loss (AddLoss)	(None,)	0	tf.math.reduce_mean_1[0][0]

Total params: 134,165  
 Trainable params: 134,165  
 Non-trainable params: 0

## Step 7: Fit model

```
In [20]: ▶ if load_weights:
    vae = vae.load_weights(args.weights)
else:
    # train the autoencoder
    vae.fit(x_train,
           epochs=epochs,
           batch_size=batch_size,
           validation_data=(x_test, None))
    vae.save_weights('vae_cnn_mnist.h5')
```

```
Epoch 1/30
313/313 [=====] - 6s 20ms/step - loss: 37.5350 - val_loss: 39.5232
Epoch 2/30
313/313 [=====] - 7s 21ms/step - loss: 37.4733 - val_loss: 39.1956
Epoch 3/30
313/313 [=====] - 7s 23ms/step - loss: 37.4376 - val_loss: 39.1108
Epoch 4/30
313/313 [=====] - 7s 21ms/step - loss: 37.3805 - val_loss: 39.0958
Epoch 5/30
313/313 [=====] - 6s 20ms/step - loss: 37.2182 - val_loss: 39.4293
Epoch 6/30
313/313 [=====] - 7s 22ms/step - loss: 37.1992 - val_loss: 39.8790
Epoch 7/30
313/313 [=====] - 7s 21ms/step - loss: 37.0978 - val_loss: 39.2028
Epoch 8/30
313/313 [=====] - 7s 22ms/step - loss: 37.0922 - val_loss: 39.0676
Epoch 9/30
313/313 [=====] - 7s 21ms/step - loss: 37.0021 - val_loss: 39.1928
Epoch 10/30
313/313 [=====] - 7s 21ms/step - loss: 36.9596 - val_loss: 39.1371
Epoch 11/30
313/313 [=====] - 6s 20ms/step - loss: 36.9170 - val_loss: 39.2005
Epoch 12/30
313/313 [=====] - 6s 20ms/step - loss: 36.8380 - val_loss: 40.0566
Epoch 13/30
313/313 [=====] - 6s 20ms/step - loss: 36.7989 - val_loss: 39.4482
Epoch 14/30
313/313 [=====] - 7s 24ms/step - loss: 36.7226 - val_loss: 39.1888
Epoch 15/30
313/313 [=====] - 7s 22ms/step - loss: 36.7134 - val_loss: 39.1914
Epoch 16/30
313/313 [=====] - 7s 22ms/step - loss: 36.6395 - val_loss: 39.0419
Epoch 17/30
313/313 [=====] - 6s 20ms/step - loss: 36.6394 - val_loss: 39.0287
Epoch 18/30
313/313 [=====] - 6s 20ms/step - loss: 36.5277 - val_loss: 39.0719
Epoch 19/30
313/313 [=====] - 6s 21ms/step - loss: 36.5322 - val_loss: 39.8637
Epoch 20/30
313/313 [=====] - 7s 21ms/step - loss: 36.4604 - val_loss: 39.8507
Epoch 21/30
313/313 [=====] - 7s 21ms/step - loss: 36.4415 - val_loss: 39.5375
Epoch 22/30
313/313 [=====] - 7s 21ms/step - loss: 36.4530 - val_loss: 39.4899
Epoch 23/30
313/313 [=====] - 6s 20ms/step - loss: 36.3539 - val_loss: 39.3501
Epoch 24/30
313/313 [=====] - 7s 21ms/step - loss: 36.3014 - val_loss: 39.6465
Epoch 25/30
313/313 [=====] - 7s 22ms/step - loss: 36.2951 - val_loss: 39.4382
Epoch 26/30
313/313 [=====] - 7s 22ms/step - loss: 36.2488 - val_loss: 39.1172
Epoch 27/30
313/313 [=====] - 8s 25ms/step - loss: 36.2677 - val_loss: 39.2752
Epoch 28/30
313/313 [=====] - 7s 21ms/step - loss: 36.1682 - val_loss: 39.2472
Epoch 29/30
313/313 [=====] - 6s 20ms/step - loss: 36.1328 - val_loss: 39.2380
Epoch 30/30
313/313 [=====] - 7s 22ms/step - loss: 36.1408 - val_loss: 39.7729
```

## Step 8: Plot new MNIST images

```
In [27]: plot_results(encoder, decoder, x_test, y_test, batch_size=batch_size)
```

